

HARITHA RAMADASS

India | haritha0849@gmail.com | +91 89408821 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL PROFILE

Dedicated and forward-thinking Data Engineer with 3.5 years of experience at Cognizant, recognised for building robust, scalable data solutions that drive real business value. Brings a strong combination of technical depth and genuine passion for the craft — approaching every project with meticulous attention to detail, engineering best practices, and a focus on solutions that are reliable and future-proof. A trusted collaborator who communicates with clarity, earns stakeholder confidence through consistent delivery, and thrives in environments that demand innovation and accountability. Eager to bring this expertise and solutions-first mindset to a high-performing team where data is treated as a competitive advantage.

TECHNICAL SKILLS

Languages: Python, SQL, PySpark

Big Data & Streaming: Apache Spark, Spark Streaming, Apache Kafka, Hadoop

Cloud: AWS (S3, Glue, Redshift, EMR, Athena, Kinesis, Lambda, IAM, CloudWatch)

Data Platforms: Databricks, Snowflake

Lakehouse & Storage: Delta Lake, Apache Iceberg, Data Lake, Data Warehouse, Medallion Architecture

ETL & Orchestration: Apache Airflow, dbt, ETL/ELT Pipelines, CDC (Change Data Capture), Incremental Loading

Data Modelling: Dimensional Modelling, Star Schema, Snowflake Schema, Fact & Dimension Tables

Data Quality & Governance: Data Quality Validation, Data Lineage, Schema Evolution, Metadata Management

Performance & Monitoring: Spark Performance Tuning, Query Optimisation, Pipeline Monitoring

DevOps & Infra: Docker, Git/GitHub

WORK EXPERIENCE

Cognizant Technology Solutions | India

Data Engineer – Associate | Sep 2022 – Aug 2025 | Healthcare clients (UK/Global)

- Engineered and maintained scalable ETL/ELT pipelines using PySpark and Apache Spark, processing 1TB+ of healthcare data weekly with consistent accuracy and reliability across downstream systems.
- Developed real-time streaming pipelines using Apache Kafka, reducing data latency by approximately 30% and improving availability across time-sensitive healthcare workflows.
- Implemented a lakehouse solution on Databricks using Delta Lake and Medallion Architecture, improving data quality and query performance across 4+ reporting datasets.
- Executed migration of on-premise workloads to AWS (S3, Glue, Redshift), contributing to a more scalable and cost-efficient cloud data environment for the project team.
- Built and managed 10+ Apache Airflow DAGs to automate pipeline scheduling and monitoring, reducing manual effort and improving workflow reliability in production.
- Developed dbt transformation models with integrated data quality checks, ensuring consistent and well-documented data for 3+ downstream consumer teams.
- Applied CDC and incremental loading strategies, reducing pipeline processing time by approximately 25% on large-scale healthcare datasets.

Data Engineer Intern | Mar 2022 – Aug 2022

- Developed and tested ETL pipelines using Python and SQL, supporting ingestion and transformation of structured healthcare data across multiple source systems.
- Executed cloud data tasks on AWS S3 and Glue, contributing to live project delivery under senior engineer guidance.
- Optimised SQL queries for data extraction and validation, improving query efficiency and ensuring accuracy across source and target systems.
- Diagnosed and resolved pipeline data quality issues, maintaining smooth and uninterrupted data operations throughout the internship period.

PROJECTS

Real-Time Banking Data Pipeline

Kafka, Debezium, AWS S3, Snowflake, dbt, Apache Airflow, Docker, Power BI

- Designed a real-time CDC pipeline using Kafka and Debezium to stream database changes into AWS S3, transformed and loaded data into Snowflake using dbt, orchestrated with Apache Airflow and visualised in Power BI.

Airbnb End-to-End Data Engineering Pipeline

AWS S3, Snowflake, dbt, Medallion Architecture, Star Schema

- Built a metadata-driven ELT pipeline ingesting raw data from AWS S3 into Snowflake, applying Medallion Architecture and dbt transformations to deliver a clean, analytics-ready Star Schema data warehouse.

Uber Real-Time Data Engineering Pipeline

Apache Spark, Azure Event Hub, Python, Star Schema

- Developed an end-to-end real-time pipeline ingesting live event streams via Azure Event Hub, processing data using Apache Spark and modelling the output into a Star Schema for analytical querying.

EDUCATION

MSc Computing: Data Analytics | Dublin City University, Ireland

2025 – 2026 | 2:1

B.Tech Computer Science Engineering | Manakula Vinayagar Institute of Technology, India

2018 – 2022 | First Class Honours | CGPA: 8.6

CERTIFICATIONS

AWS Certified Data Engineer

AWS Certified Solutions Architect

AWS Certified Cloud Practitioner